

Verify Docker installation.

In this setup we are running Docker version 29.1.3, but any recent stable release works.

```
tester@vbox:~$ docker --version
Docker version 29.1.3, build f52814d
tester@vbox:~$ kubectl version --client
Client Version: v1.35.0
Kustomize Version: v5.7.1
tester@vbox:~$ helm version
version.BuildInfo{Version:"v3.20.0", GitCommit:"b2e4314fa0f229a1de7b4c981273f61d69ee5a59", GitTreeState:"clean", GoVersion:"go1.25.6"}
```

Set the required environment variables for PostgreSQL and the Backend.

The repository includes a `.env.example` file that can be copied and used as a template for local development.

```
● tester@vbox:~/wellnes-ops$ cp .env.example env/dev/.env
● tester@vbox:~/wellnes-ops$ cat env/dev/.env
# =====
# PostgreSQL (OBLIGATORIO)
# =====
POSTGRES_DB=wellness
POSTGRES_USER=postgres
POSTGRES_PASSWORD=wellness

tester@vbox:~$ git clone https://github.com/luisrodvilladaorg/wellnes-ops.git
cd wellnes-ops
Cloning into 'wellnes-ops'...
remote: Enumerating objects: 9178, done.
remote: Counting objects: 100% (58/58), done.
remote: Compressing objects: 100% (41/41), done.
remote: Total 9178 (delta 28), reused 34 (delta 17), pack-reused 9120 (from 3)
Receiving objects: 100% (9178/9178), 21.78 MiB | 3.45 MiB/s, done.
Resolving deltas: 100% (2380/2380), done.
tester@vbox:~/wellnes-ops$
```

Start the container stack using Docker Compose.

In this example we use the `dev` environment, but the `prod` environment can be used the same way depending on the deployment scenario.

```
docker compose -f docker-compose-dev up -d --build
```

```
tester@vbox:~/wellnes-ops$ docker compose -f docker-compose.dev.yml up -d --build
[+] Building 0.6s (28/28) FINISHED
=> [internal] load local bake definitions
=> => reading from stdin 1.42kB
=> [frontend internal] load build definition from Dockerfile.dev
=> => transferring dockerfile: 583B
=> [nginx internal] load build definition from Dockerfile.dev
=> => transferring dockerfile: 284B
=> [backend internal] load build definition from Dockerfile.dev
=> => transferring dockerfile: 364B
=> [nginx internal] load metadata for docker.io/library/nginx:stable-alpine
=> [backend internal] load metadata for docker.io/library/node:20-alpine
=> [frontend internal] load .dockerignore
=> => transferring context: 2B
=> [nginx internal] load .dockerignore
=> => transferring context: 2B
=> [frontend 1/3] FROM docker.io/library/nginx:stable-alpine@sha256:bb2ba4172e705075aca7f0f51a21fd7c758e543e09b220a4eba5a067666180a5
=> => resolve docker.io/library/nginx:stable-alpine@sha256:bb2ba4172e705075aca7f0f51a21fd7c758e543e09b220a4eba5a067666180a5
=> [frontend internal] load build context
```

For security reasons, the Backend service is not exposed externally.

To validate connectivity, exec into the Backend container and test the internal endpoint directly.

```
docker exec -it wellness-backend-container sh  
wget -qO- http://localhost:3000/api/health
```

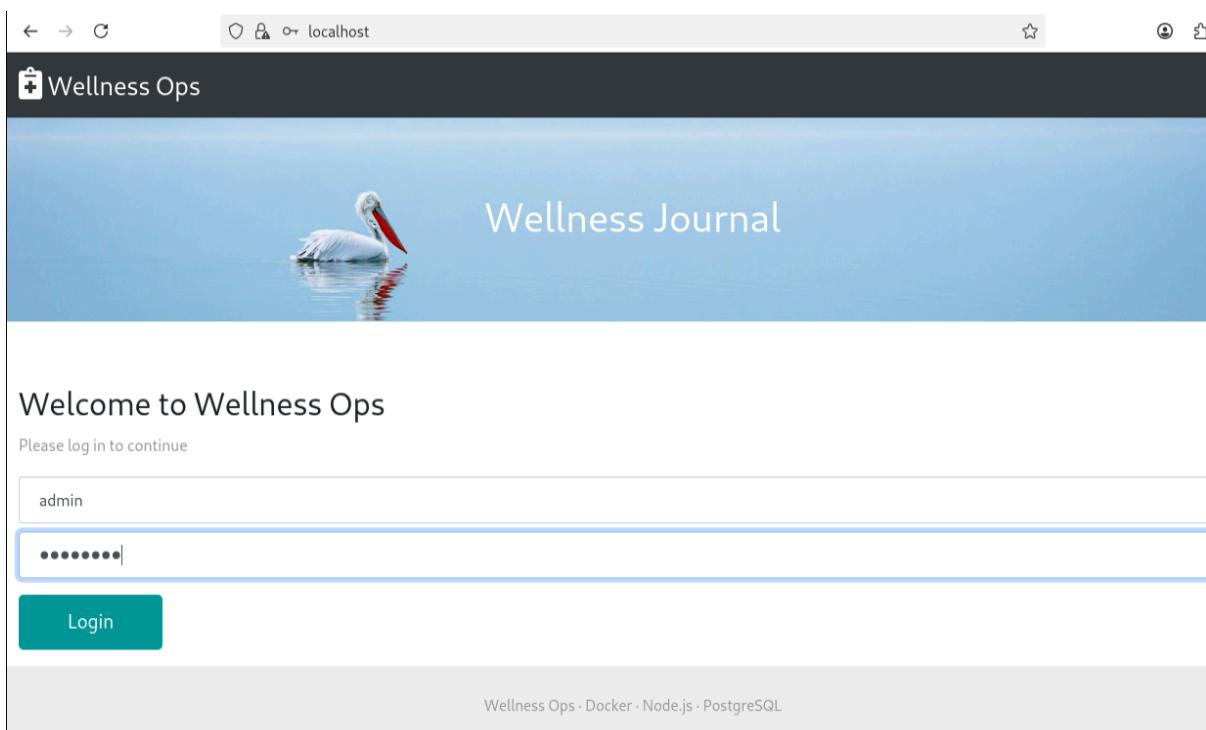
```
tester@vbox:~/wellnes-ops$ docker ps  
CONTAINER ID IMAGE COMMAND CREATED STATUS PORTS  
NAMES  
fc68a70799ee wellnes-ops-nginx "/docker-entrypoint..." About a minute ago Up About a minute 0.0.0.0:80->80/tcp, [::]:80->80/tcp, 0.0.0.0:443->443/tcp, [::]:443->443/tcp  
tcp wellnes-nginx-proxy  
80e44fb919ed wellnes-ops-frontend "/docker-entrypoint..." About a minute ago Up About a minute (healthy) 0.0.0.0:8080->80/tcp, [::]:8080->80/tcp  
wellness-frontend-container  
88a7158d8e4 wellnes-ops-backend "/docker-entrypoint..." About a minute ago Up About a minute (healthy) 3000/tcp  
wellness-backend-container  
7bc0de7f4170 grafana/grafana:latest "/run.sh"  
wellness-grafana  
ea964b74bcfd prom/prometheus:latest "/bin/prometheus --c..."  
counts - Sign in requested ethus  
ed4d1063d8b3 postgres:15-alpine "docker-entrypoint..." About a minute ago Up About a minute (healthy) 5432/tcp  
wellness-postgres-db
```

```
tester@vbox:~/wellnes-ops$ docker exec -it wellness-backend-container sh  
/app # wget -qO- http://localhost:3000/api/health  
{"status": "OK"} /app #  
/app # █
```

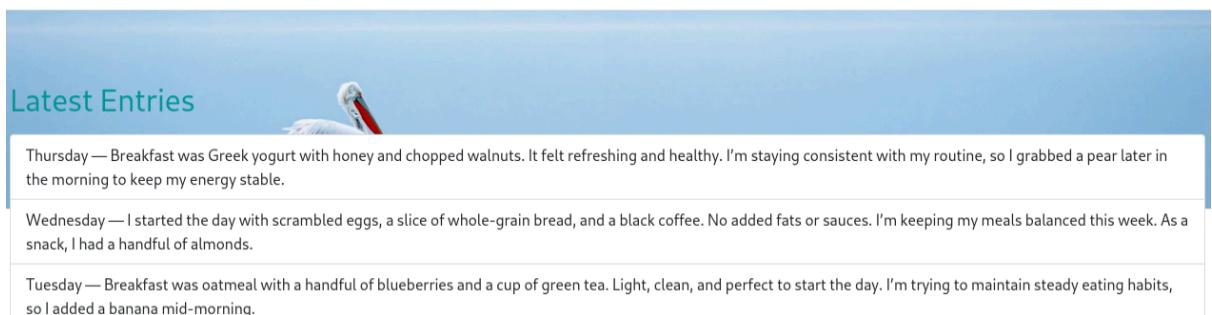
Access the application through any web browser using TLS:

<https://localhost>

Log in using the default development credentials: admin / admin123



Create and save a few records to validate that data persistence works as expected.



The screenshot shows a web application titled "Latest Entries". It displays three meal entries:

- Thursday — Breakfast was Greek yogurt with honey and chopped walnuts. It felt refreshing and healthy. I'm staying consistent with my routine, so I grabbed a pear later in the morning to keep my energy stable.
- Wednesday — I started the day with scrambled eggs, a slice of whole-grain bread, and a black coffee. No added fats or sauces. I'm keeping my meals balanced this week. As a snack, I had a handful of almonds.
- Tuesday — Breakfast was oatmeal with a handful of blueberries and a cup of green tea. Light, clean, and perfect to start the day. I'm trying to maintain steady eating habits, so I added a banana mid-morning.

Add New Entry

Title

Description

Once the application is confirmed to be running correctly, shut it down before proceeding with the next steps..

```
docker compose -f docker-compose.dev.yml down
```

```
● tester@vbox:~/wellnes-ops$ docker compose -f docker-compose.dev.yml down
[+] Running 8/8
✓ Container wellness-grafana           Removed
✓ Container wellness-nginx-proxy       Removed
✓ Container wellness-prometheus        Removed
✓ Container wellness-frontend-container Removed
✓ Container wellness-backend-container  Removed
✓ Container wellness-postgres-db       Removed
✓ Network wellnes-ops_web_net         Removed
✓ Network wellnes-ops_backend_net     Removed
○ tester@vbox:~/wellnes-ops$
```

KUBERNETES

Verify kubectl Installation

Ensure that `kubectl` is installed and accessible from the system. This confirms that the Kubernetes CLI is available before interacting with the cluster.

```
tester@vbox:~/wellnes-ops$ kubectl version
Client Version: v1.35.0
Kustomize Version: v5.7.1
The connection to the server localhost:8080 was refused - did you specify the right host or port?
tester@vbox:~/wellnes-ops$
```

Create the Kubernetes cluster

For local environments—such as this setup running Debian 12 inside a VirtualBox VM—we provision a Kubernetes cluster. In this example, the cluster is named `cluster-wellness-local`

Run the following commands to initialize the cluster

```
k3d cluster create cluster-wellness-local \
--api-port 6443 \
--servers 1 \
--agents 0 \
--port 80:80@loadbalancer \
--port 443:443@loadbalancer \
--k3s-arg "--disable=traefik@server:0"
```

```
tester@vbox:~/wellnes-ops$ k3d cluster create cluster-wellness-local \
--api-port 6443 \
--servers 1 \
--agents 0 \
--port 80:80@loadbalancer \
--port 443:443@loadbalancer \
--k3s-arg "--disable=traefik@server:0"
INFO[0000] portmapping '80:80' targets the loadbalancer: defaulting to [servers:*:proxy agents:*:proxy]
INFO[0000] portmapping '443:443' targets the loadbalancer: defaulting to [servers:*:proxy agents:*:proxy]
INFO[0000] Prep: Network
INFO[0000] Created network 'k3d-cluster-wellness-local'
INFO[0000] Created image volume k3d-cluster-wellness-local-images
INFO[0000] Starting new tools node...
INFO[0000] Starting node 'k3d-cluster-wellness-local-tools'
INFO[0001] Creating node 'k3d-cluster-wellness-local-server-0'
```

Step by step

Create the Kubernetes configuration directory if it does not already exist:

Extract the cluster kubeconfig and write it to the default configuration file: This ensures that kubectl interacts with the correct cluster context.

Set the appropriate file permissions: Restrict access so only the current user can read the kubeconfig

List available contexts to verify that the configuration was applied correctly:

Commands:

```
mkdir -p ~/.kube
```

```
k3d kubeconfig get cluster-wellness-local > ~/.kube/config
```

```
chmod 600 ~/.kube/config
```

```
kubectl config get-contexts
```

```
● tester@vbox:~/wellnes-ops$ mkdir -p ~/.kube
● tester@vbox:~/wellnes-ops$ k3d kubeconfig get cluster-wellness-local > ~/.kube/config
● tester@vbox:~/wellnes-ops$ chmod 600 ~/.kube/config
● tester@vbox:~/wellnes-ops$ kubectl config get-contexts
  CURRENT   NAME                 CLUSTER          AUTHINFO           NAMESPACE
  *         k3d-cluster-wellness-local   k3d-cluster-wellness-local  [REDACTED]@k3d-cluster-wellness-local
● tester@vbox:~/wellnes-ops$ kubectl config use-context k3d-cluster-wellness-local
Switched to context "k3d-cluster-wellness-local".
● tester@vbox:~/wellnes-ops$ kubectl get nodes
  NAME                  STATUS    ROLES           AGE     VERSION
  k3d-cluster-wellness-local-server-0   Ready    control-plane,master   7m17s   v1.31.5+k3s1
● tester@vbox:~/wellnes-ops$
```

Verify that the Kubernetes cluster is installed and running.

Since this environment uses k3d, confirm that the cluster is active and recognized by the local system

Check that the Kubernetes nodes are registered and in a Ready state

This ensures the control plane and worker nodes are operational.

k3d cluster list

kubectl get nodes

```
● tester@vbox:~/wellnes-ops$ k3d cluster list
  NAME          SERVERS   AGENTS   LOADBALANCER
cluster-wellness-local  1/1      0/0      true
● tester@vbox:~/wellnes-ops$ kubectl get nodes
  NAME           STATUS    ROLES      AGE   VERSION
k3d-cluster-wellness-local-server-0   Ready     control-plane,master   96s   v1.31.5+k3s1
○ tester@vbox:~/wellnes-ops$ █
```

List the pods running in the `kube-system` namespace.

This allows you to verify that core Kubernetes components (DNS, CNI, metrics, controllers, etc.) are healthy and running as expected.

kubectl get pods -n kube-system

```
● tester@vbox:~/wellnes-ops$ kubectl get pods -n kube-system
  NAME          READY   STATUS    RESTARTS   AGE
coredns-ccb9694c-g5rs6   1/1     Running   0          2m40s
local-path-provisioner-5cf85fd84d-5vnpm   1/1     Running   0          2m40s
metrics-server-5985cbc9d7-xt4z4   1/1     Running   0          2m40s
● tester@vbox:~/wellnes-ops$ kubectl apply -f https://raw.githubusercontent.com/kubernetes/ingress-nginx/controller-v1.11.1/deploy/static/provider/cloud/deployment.yaml
namespace/ingress-nginx created
serviceaccount/ingress-nginx created
serviceaccount/ingress-nginx-admission created
role.rbac.authorization.k8s.io/ingress-nginx created
role.rbac.authorization.k8s.io/ingress-nginx-admission created
clusterrole.rbac.authorization.k8s.io/ingress-nginx created
clusterrole.rbac.authorization.k8s.io/ingress-nginx-admission created
rolebinding.rbac.authorization.k8s.io/ingress-nginx created
rolebinding.rbac.authorization.k8s.io/ingress-nginx-admission created
```

Install Required Tools for the Application

Once the cluster is up and running, install the required tooling for the application stack. This includes Helm, the NGINX Ingress Controller (we replace the default Traefik shipped with k3d), MetallB, and the MetallB IP address pool. Finally, apply the necessary Kubernetes manifests.

Install Helm using the official installation script:

```
curl https://raw.githubusercontent.com/helm/helm/main/scripts/get-helm-3 | bash
```

Verify the installation:

```
● tester@vbox:~/wellnes-ops$ curl https://raw.githubusercontent.com/helm/helm/main/scripts/get-helm-3 | bash
  % Total    % Received % Xferd  Average Speed   Time     Time   Current
          Dload  Upload   Total Spent   Left  Speed
100 11929 100 11929    0      0  114k      0 --:--:--:--:--:--:--:-- 115k
Helm v3.20.0 is already latest
● tester@vbox:~/wellnes-ops$ helm version
version.BuildInfo{Version:"v3.20.0", GitCommit:"b2e4314fa0f229a1de7b4c981273f61d69ee5a59", GitTreeState:"clean", GoVersion:"go1.25.6"}
```

Add the official NGINX Ingress Controller Helm repository and update the index:

```
helm repo add ingress-nginx https://kubernetes.github.io/ingress-nginx
helm repo update
```

Install the NGINX Ingress Controller:

This deployment replaces the default Traefik controller used by k3d and enables NGINX-based ingress routing. The flag

`controller.watchIngressWithoutClass=true` allows NGINX to manage ingresses that do not explicitly define an `ingressClass`

```
helm install ingress-nginx ingress-nginx/ingress-nginx \
--namespace ingress-nginx \
--create-namespace \
--set controller.watchIngressWithoutClass=true
```

```
● tester@vbox:~/wellnes-ops$ helm install ingress-nginx ingress-nginx/ingress-nginx \
--namespace ingress-nginx \
--create-namespace \
--set controller.watchIngressWithoutClass=true
NAME: ingress-nginx
LAST DEPLOYED: Mon Feb  9 12:48:43 2026
NAMESPACE: ingress-nginx
STATUS: deployed
REVISION: 1
TEST SUITE: None
NOTES:
The ingress-nginx controller has been installed.
It may take a few minutes for the load balancer IP to be available.
You can watch the status by running 'kubectl get service --namespace ingress-nginx ingress-nginx-controller --output wide --watch'

An example Ingress that makes use of the controller:
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: example
  namespace: foo
spec:
  ingressClassName: nginx
```

Validate NGINX Installation and Deploy MetalLB

This ensures that the ingress layer is operational before deploying any load-balancing components.

```
kubectl get pods -n ingress-nginx
```

```
● tester@vbox:~/wellnes-ops$ kubectl get pods -n ingress-nginx
kubectl get svc -n ingress-nginx
NAME                                READY   STATUS    RESTARTS   AGE
ingress-nginx-controller-56bbb78745-m77s8   1/1     Running   0          77s
NAME                            TYPE      CLUSTER-IP        EXTERNAL-IP   PORT(S)           AGE
ingress-nginx-controller   LoadBalancer   10.43.116.178   172.18.0.2   80:31351/TCP,443:30834/TCP   77s
ingress-nginx-controller-admission ClusterIP   10.43.24.4       <none>        443/TCP          77s
```

Install MetalLB using the official manifest.

MetalLB provides Layer 2 load balancing for bare-metal and local Kubernetes clusters, including k3d environments.

```
kubectl apply -f
```

```
https://raw.githubusercontent.com/metallb/metallb/v0.14.5/config/manifests/metallb-native.yaml
```

```
● tester@vbox:~/wellnes-ops$ kubectl apply -f https://raw.githubusercontent.com/metallb/metallb/v0.14.5/config/manifests/metallb-native.yaml
namespace/metallb-system created
customresourcedefinition.apiextensions.k8s.io/bfdprofiles.metallb.io created
customresourcedefinition.apiextensions.k8s.io/bgpadvertisements.metallb.io created
customresourcedefinition.apiextensions.k8s.io/bgppeers.metallb.io created
customresourcedefinition.apiextensions.k8s.io/communities.metallb.io created
customresourcedefinition.apiextensions.k8s.io/ipaddresspools.metallb.io created
customresourcedefinition.apiextensions.k8s.io/l2advertisements.metallb.io created
customresourcedefinition.apiextensions.k8s.io/serviceclustatuses.metallb.io created
serviceaccount/controller created
serviceaccount/speaker created
role.rbac.authorization.k8s.io/controller created
role.rbac.authorization.k8s.io/pod-lister created
clusterrole.rbac.authorization.k8s.io/metallb-system:controller created
clusterrole.rbac.authorization.k8s.io/metallb-system:speaker created
rolebinding.rbac.authorization.k8s.io/controller created
rolebinding.rbac.authorization.k8s.io/pod-lister created
clusterrolebinding.rbac.authorization.k8s.io/metallb-system:controller created
clusterrolebinding.rbac.authorization.k8s.io/metallb-system:speaker created
configmap/metallb-exclude12 created
secret/metallb-webhook-cert created
```

Validate MetalLB and Install IP Address Pool + Prometheus

Verify that MetalLB pods are running correctly:

This confirms that the load balancer components are healthy before configuring the address pool.

```
kubectl get pods -n metallb-system
```

```
● tester@vbox:~/wellnes-ops$ kubectl get pods -n metallb-system
  NAME           READY   STATUS    RESTARTS   AGE
  controller-58859b4d4f-zn4qc   1/1     Running   0          53s
  speaker-wzqtw      1/1     Running   0          53s
● tester@vbox:~/wellnes-ops$
```

Install the MetalLB IP address pool configuration: This defines the range of IPs that MetalLB can allocate for LoadBalancer services.

Validate that the IPAddressPool resource has been created:

```
kubectl apply -f /k8s/metallb/
```

```
kubectl get ipaddresspools -n metallb-system
```

```
● tester@vbox:~/wellnes-ops$ kubectl apply -f k8s/metallb/
l2advertisement.metallb.io/local-l2 created
ipaddresspool.metallb.io/local-pool created
● tester@vbox:~/wellnes-ops$ kubectl get ipaddresspools -n metallb-system
  NAME      AUTO ASSIGN   AVOID BUGGY IPS   ADDRESSES
  local-pool  true        false            ["172.19.255.200-172.19.255.250"]
```

Install Prometheus in an isolated namespace

Add the official Prometheus Helm repository and update the index

Create a dedicated namespace for monitoring components

Install the latest Prometheus chart

`kubectl get pods -n monitoring`

Commands:

```
helm repo add prometheus-community https://prometheus-community.github.io/helm-charts
helm repo update
```

`kubectl create namespace monitoring`

`helm install prometheus prometheus-community/prometheus -n monitoring`

`kubectl get pods -n monitoring`

```
● tester@vbox:~/wellnes-ops$ kubectl get pods -n monitoring
  NAME                               READY   STATUS    RESTARTS   AGE
  prometheus-alertmanager-0          1/1     Running   0          97s
  prometheus-kube-state-metrics-74bb7cc548-f722f   1/1     Running   0          97s
  prometheus-prometheus-node-exporter-slln7        1/1     Running   0          97s
  prometheus-prometheus-pushgateway-665f7ffd4c-5qgls 1/1     Running   0          97s
  prometheus-server-6496dc96-wvtx                2/2     Running   0          97s
○ tester@vbox:~/wellnes-ops$
```

Install Required CRDs

Install the required Custom Resource Definitions (CRDs). These CRDs are needed by the Prometheus stack and must be applied before deploying additional monitoring components.

helm repo add prometheus-community <https://prometheus-community.github.io/helm-charts>

```
● tester@vbox:~/wellnes-ops$ helm repo add prometheus-community https://prometheus-community.github.io/helm-charts
helm repo update
"prometheus-community" has been added to your repositories
Hang tight while we grab the latest from your chart repositories...
...Successfully got an update from the "ingress-nginx" chart repository
...Successfully got an update from the "prometheus-community" chart repository
Update Complete. *Happy Helm-ing!*
● tester@vbox:~/wellnes-ops$ kubectl create namespace monitoring
namespace/monitoring created
● tester@vbox:~/wellnes-ops$ helm install prometheus prometheus-community/prometheus -n monitoring
NAME: prometheus
LAST DEPLOYED: Tue Feb 10 11:04:58 2026
NAMESPACE: monitoring
STATUS: deployed
REVISION: 1
TEST SUITE: None
NOTES:
The Prometheus server can be accessed via port 80 on the following DNS name from within your cluster:
prometheus-server.monitoring.svc.cluster.local
```

Verify that the ServiceMonitor CRD is installed: This ensures that the monitoring stack can register scrape targets correctly.

kubectl get crds | grep servicemonitor

```
● tester@vbox:~/wellnes-ops$ kubectl get crds | grep servicemonitor
  servicemonitors.monitoring.coreos.com          2026-02-10T10:17:30Z
○ tester@vbox:~/wellnes-ops$ █
```

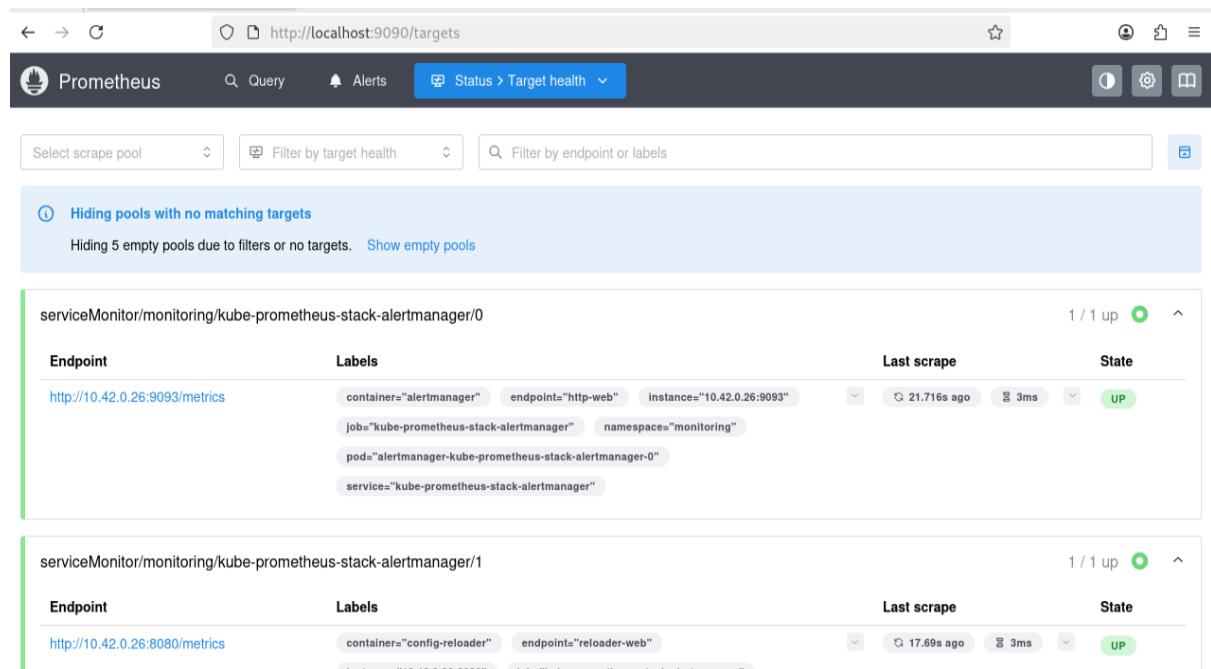
Apply Monitoring Manifests and Expose Prometheus

kubectl apply -f k8s/monitoring/

Expose the Prometheus UI locally using port-forwarding: This forwards traffic from port 9090 on your machine to the Prometheus pod.

```
● tester@vbox:~/wellnes-ops$ kubectl apply -f k8s/monitoring/
servicemonitor.monitoring.coreos.com/backend-monitoring unchanged
○ tester@vbox:~/wellnes-ops$ kubectl port-forward -n monitoring pod/prometheus-kube-prometheus-stack-prometheus-0 9090:9090
Forwarding from 127.0.0.1:9090 -> 9090
Forwarding from [::1]:9090 -> 9090
█
```

Validate access and confirm that Prometheus is running (status: Up). Open your browser and navigate to: <http://localhost:9090>



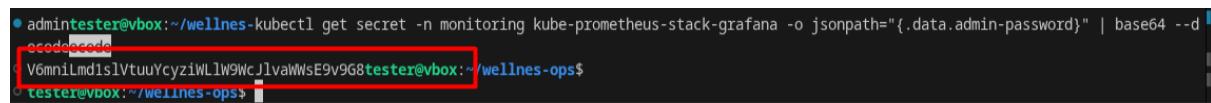
The screenshot shows the Prometheus web interface at <http://localhost:9090/targets>. The top navigation bar includes links for Query, Alerts, and Status > Target health. The main content area displays two serviceMonitor entries under the monitoring namespace:

- serviceMonitor/monitoring/kube-prometheus-stack-alertmanager/0**: Shows 1 target up. The target is <http://10.42.0.26:9093/metrics>. Labels include container=alertmanager, endpoint=http-web, instance=10.42.0.26:9093, job=kube-prometheus-stack-alertmanager, namespace=monitoring, pod=alertmanager-kube-prometheus-stack-alertmanager-0, and service=kube-prometheus-stack-alertmanager.
- serviceMonitor/monitoring/kube-prometheus-stack-alertmanager/1**: Shows 1 target up. The target is <http://10.42.0.26:8080/metrics>. Labels include container=config-reloader, endpoint=reloader-web, instance=10.42.0.26:8080, and job=kube-prometheus-stack-alertmanager.

Grafana is included by default in the Prometheus stack

To access the dashboard, retrieve the auto-generated admin password from the Kubernetes secret:

```
kubectl get secret -n monitoring kube-prometheus-stack-grafana -o jsonpath='{.data.admin-password}' | base64 --decode
```



```
● admin[tester@vbox:~/wellnes] kubectl get secret -n monitoring kube-prometheus-stack-grafana -o jsonpath='{.data.admin-password}' | base64 --decode
[REDACTED]
V6mnILmd1s1VtuoYcyziWLlW9WcJlvaWWsE9v9G8[REDACTED]
tester@vbox:~/wellnes$
```

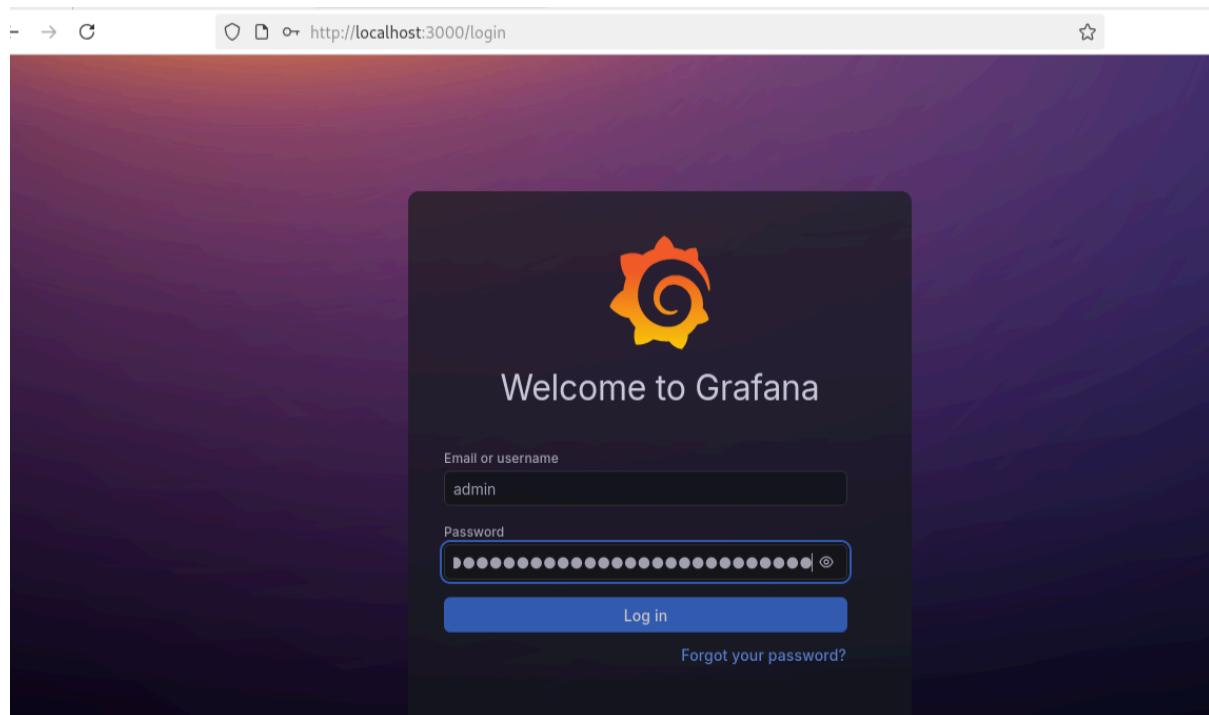
Expose the Grafana dashboard locally using port-forwarding. This allows access to the UI without exposing the service externally.

```
kubectl port-forward -n monitoring deploy/kube-prometheus-stack-grafana 3000:3000
```

```
tester@vbox:~/wellnes-ops$ kubectl port-forward -n monitoring deploy/kube-prometheus-stack-grafana 3000:3000
Forwarding from 127.0.0.1:3000 -> 3000
Forwarding from [::1]:3000 -> 3000
```

Access the Grafana dashboard through the local port-forwarding session. Once the port-forward is active, open your browser and navigate to:

<http://localhost:3000>



Install Cert-Manager

To enable TLS certificate management within the cluster, install Cert-Manager and configure the required certificate resources. Cert-Manager automates the issuance and renewal of TLS certificates for Kubernetes workloads

```
kubectl apply -f
```

```
https://github.com/cert-manager/cert-manager/releases/download/v1.15.1/cert-manager.yaml
```

```
● tester@vbox:~/wellnes-ops$ kubectl apply -f https://github.com/cert-manager/cert-manager/releases/download/v1.15.1/cert-manager.yaml
namespace/cert-manager created
customresourcedefinition.apirextensions.k8s.io/certificaterequests.cert-manager.io created
customresourcedefinition.apirextensions.k8s.io/certificates.cert-manager.io created
customresourcedefinition.apirextensions.k8s.io/challenges.acme.cert-manager.io created
customresourcedefinition.apirextensions.k8s.io/clusterissuers.cert-manager.io created
customresourcedefinition.apirextensions.k8s.io/issuers.cert-manager.io created
customresourcedefinition.apirextensions.k8s.io/orders.acme.cert-manager.io created
serviceaccount/cert-manager-cainjector created
serviceaccount/cert-manager created
serviceaccount/cert-manager-webhook created
clusterrole.rbac.authorization.k8s.io/cert-manager-cainjector created
clusterrole.rbac.authorization.k8s.io/cert-manager-controller-issuers created
clusterrole.rbac.authorization.k8s.io/cert-manager-controller-clusterissuers created
clusterrole.rbac.authorization.k8s.io/cert-manager-controller-certificates created
clusterrole.rbac.authorization.k8s.io/cert-manager-controller-orders created
clusterrole.rbac.authorization.k8s.io/cert-manager-controller-challenges created
clusterrole.rbac.authorization.k8s.io/cert-manager-controller-ingress-shim created
clusterrole.rbac.authorization.k8s.io/cert-manager-cluster-view created
clusterrole.rbac.authorization.k8s.io/cert-manager-view created
clusterrole.rbac.authorization.k8s.io/cert-manager-edit created
```

Confirm Cert-Manager Pods Are Running

Verify that the Cert-Manager components are installed and running correctly.

This ensures that the certificate management system is fully operational before creating Issuers or Certificates.

```
● tester@vbox:~/wellnes-ops$ kubectl get pods -n cert-manager
NAME                               READY   STATUS    RESTARTS   AGE
cert-manager-675d667c9-7vmcc       1/1     Running   0          46s
cert-manager-cainjector-6674494d8-cskvf 1/1     Running   0          46s
cert-manager-webhook-8566bc98-2mx7p   1/1     Running   0          46s
```

Apply the application manifests one by one, in the correct order.
It is important to wait ~30 seconds between each deployment to ensure that the pods reach a Running state.

This is required because:

- The backend depends on the database.
- The NGINX ingress depends on the frontend even if they run in isolated namespaces.

```
kubectl apply -f /k8s/postgres/  
kubectl apply -f /k8s/backend/  
kubectl apply -f /k8s/frontend/  
kubectl apply -f /k8s/nginx/
```

```
● tester@vbox:~/wellnes-ops$ kubectl apply -f k8s/postgres/  
configmap/postgres-init created  
job.batch/postgres-init created  
secret/postgres-secret created  
service/postgres-service created  
statefulset.apps/postgres created  
● tester@vbox:~/wellnes-ops$ kubectl get pods  
NAME READY STATUS RESTARTS AGE  
postgres-0 1/1 Running 0 36s  
postgres-init-ntk7s 0/1 Completed 0 36s  
● tester@vbox:~/wellnes-ops$ █
```

```
● tester@vbox:~/wellnes-ops$ kubectl apply -f k8s/backend/  
configmap/backend-config created  
deployment.apps/backend created  
secret/backend-jwt-secret created  
secret/backend-secret created  
service/backend created  
● tester@vbox:~/wellnes-ops$ kubectl get pods  
NAME READY STATUS RESTARTS AGE  
backend-5d75cc8db8-xr5fh 0/1 PodInitializing 0 15s  
postgres-0 1/1 Running 0 10m  
postgres-init-ntk7s 0/1 Completed 0 10m  
○ tester@vbox:~/wellnes-ops$ █
```

```
● tester@vbox:~/wellnes-ops$ kubectl apply -f k8s/frontend/  
deployment.apps/frontend created  
service/frontend-service created  
○ tester@vbox:~/wellnes-ops$ █
```

```
● tester@vbox:~/wellnes-ops$ kubectl apply -f k8s/nginx/
configmap/nginx-gateway-config created
deployment.apps/nginx-gateway created
service/nginx-gateway created
○ tester@vbox:~/wellnes-ops$
```

Verify PostgreSQL Installation as a StatefulSet

Confirm that PostgreSQL is deployed as a StatefulSet (not just as a Pod). This ensures that the database is using persistent identity and stable storage, as expected for stateful workloads.

```
● usuario1@vbox:/opt/wellnes-ops$ kubectl get statefulsets
  NAME      READY   AGE
  postgres   1/1     2d3h
○ usuario1@vbox:/opt/wellnes-ops$
```

Verify That All Components Are Running

Confirm that all deployed components are in a Running state.

This ensures that every part of the application stack—database, backend, frontend, and NGINX—is healthy and ready to serve traffic.

```
● tester@vbox:~/wellnes-ops$ kubectl get pods
  NAME                  READY   STATUS    RESTARTS   AGE
  backend-5d75cc8db8-xr5fh   1/1     Running   0          82s
  frontend-6f9fd8c7c6-zqj9l   1/1     Running   0          11s
  nginx-gateway-769d76746d-dwvgw   1/1     Running   2 (26s ago)   33s
  postgres-0                1/1     Running   0          11m
  postgres-init-ntk7s        0/1     Completed  0          11m
```

Install and Verify Cert-Manager

Install and configure Cert-Manager.

This component is required for automated TLS certificate issuance and renewal inside the cluster.

```
kubectl create namespace cert-manager  
helm repo add jetstack https://charts.jetstack.io  
helm repo update
```

Install Cert-Manager with CRDs enabled:

```
helm install cert-manager jetstack/cert-manager \  
--namespace cert-manager \  
--set crds.enabled=true
```

Verify that Cert-Manager pods are running:

```
kubectl get pods -n cert-manager
```

List all resources in the Cert-Manager namespace:

```
kubectl get all -n cert-manager
```

```
● tester@vbox:~/wellnes-ops$ kubectl get pods -n cert-manager  
kubectl get all -n cert-manager  
NAME                                     READY   STATUS    RESTARTS   AGE  
cert-manager-675d667c9-sgpfk              1/1     Running   1 (9m1s ago)  64m  
cert-manager-cainjector-6674494d8-dfkm4   1/1     Running   1 (9m1s ago)  64m  
cert-manager-webhook-8566bcbc98-rk28g     1/1     Running   1 (9m1s ago)  64m  
  
NAME                                     READY   STATUS    RESTARTS   AGE  
pod/cert-manager-675d667c9-sgpfk          1/1     Running   1 (9m1s ago)  64m  
pod/cert-manager-cainjector-6674494d8-dfkm4 1/1     Running   1 (9m1s ago)  64m  
pod/cert-manager-webhook-8566bcbc98-rk28g  1/1     Running   1 (9m1s ago)  64m  
  
NAME           TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE  
service/cert-manager  ClusterIP  10.43.17.125  <none>        9402/TCP   64m  
service/cert-manager-webhook  ClusterIP  10.43.134.240 <none>        443/TCP    64m  
  
NAME           READY   UP-TO-DATE   AVAILABLE   AGE  
deployment.apps/cert-manager            1/1     1           1           64m  
deployment.apps/cert-manager-cainjector 1/1     1           1           64m  
deployment.apps/cert-manager-webhook    1/1     1           1           64m  
  
NAME           DESIRED  CURRENT   READY   AGE  
replicaset.apps/cert-manager-675d667c9  1        1        1       64m  
replicaset.apps/cert-manager-cainjector-6674494d8 1        1        1       64m  
replicaset.apps/cert-manager-webhook-8566bcbc98  1        1        1       64m
```

Apply TLS Manifests (In Order)

Apply the TLS configuration manifests located in `/k8s/tls/`, one by one and in the correct order. This sequencing is important to ensure proper initialization of Issuers, Certificates, and all TLS-related resources required for HTTPS

```
kubectl apply -f k8s/tls/ca-clusterissuer.yml
```

```
kubectl describe clusterissuer wellness-ca
```

```
● tester@vbox:~/wellnes-ops$ kubectl apply -f k8s/tls/ca-clusterissuer.yml
clusterissuer.cert-manager.io/wellness-ca unchanged
● tester@vbox:~/wellnes-ops$ kubectl describe clusterissuer wellness-ca
Name:          wellness-ca
Namespace:    
Labels:        <none>
Annotations:   <none>
API Version:  cert-manager.io/v1
Kind:          ClusterIssuer
Metadata:
  Creation Timestamp:  2026-02-10T12:18:27Z
  Generation:        1
  Resource Version:  12090
  UID:              5a8094da-ef17-46b1-a7fe-e723ddaa2989f
Spec:
  Ca:
    Secret Name:  wellness-ca-secret
Status:
  Conditions:
    Last Transition Time:  2026-02-10T12:43:32Z
    Message:        Signing ClusterIssuer
    Status:        True
    Type:          SigningClusterIssuer
```

Apply the CA-signed Certificate manifest:

```
kubectl apply -f k8s/tls/ca-certificate.yml
```

```
kubectl get certificate -n cert-manager
```

```
kubectl get certificate -n cert-manager
```

```
● tester@vbox:~/wellnes-ops$ kubectl apply -f k8s/tls/ca-certificate.yml
certificate.cert-manager.io/wellness-ca unchanged
● tester@vbox:~/wellnes-ops$ kubectl get certificate -n cert-manager
NAME      READY   SECRET           AGE
wellness-ca  True    wellness-ca-secret  17m
```

Apply the TLS certificate manifest:

```
kubectl apply -f k8s/tls/wellness-tls.yml
```

Verify that the TLS secret has been created in the default namespace:

```
kubectl get secret wellness-tls -n default
```

```
● tester@vbox:~/wellnes-ops$ kubectl apply -f k8s/tls/wellness-tls.yml
certificate.cert-manager.io/wellness-tls unchanged
● tester@vbox:~/wellnes-ops$ kubectl get secret wellness-tls -n default
NAME          TYPE        DATA   AGE
wellness-tls  kubernetes.io/tls  3      15m
```

Apply the Ingress manifest for the application:

```
kubectl apply -f k8s/tls/wellness-ingress.yml
```

Verify that the Ingress resource has been created and is configured correctly:

```
kubectl describe ingress wellness-ingress -n default
```

```
● tester@vbox:~/wellnes-ops$ kubectl apply -f k8s/tls/wellness-ingress.yml
ingress.networking.k8s.io/wellness-ingress unchanged
● tester@vbox:~/wellnes-ops$ kubectl describe ingress wellness-ingress -n default
Name:            wellness-ingress
Labels:          <none>
Namespace:       default
Address:         172.19.255.200
Ingress Class:  nginx
Default backend: <default>
TLS:
  wellness-tls terminates wellness.local
Rules:
  Host           Path  Backends
  ----          ----  -----
  wellness.local
    /api      backend:3000 (10.42.0.93:3000)
    /        frontend-service:80 (10.42.0.82:80)
Annotations:    cert-manager.io/cluster-issuer: wellness-ca
                nginx.ingress.kubernetes.io/force-ssl-redirect: true
                nginx.ingress.kubernetes.io/ssl-redirect: true
Events:
```

Verify Application Access

Confirm that the application is accessible.

Once all components are deployed and all TLS resources are applied, validate that the application is reachable through its HTTPS endpoint.

Typical checks include:

- Accessing the application URL in a browser
- Confirming that the TLS certificate is valid
- Ensuring that the Ingress routes traffic correctly
- Verifying that the backend responds without errors

```
● tester@vbox:~/wellnes-ops$ curl -vk https://wellness.local
*   Trying 127.0.0.1:443...
* Connected to wellness.local (127.0.0.1) port 443 (#0)
* ALPN: offers h2,http/1.1
* TLSv1.3 (OUT), TLS handshake, Client hello (1):
* TLSv1.3 (IN), TLS handshake, Server hello (2):
* TLSv1.3 (IN), TLS handshake, Encrypted Extensions (8):
* TLSv1.3 (IN), TLS handshake, Certificate (11):
* TLSv1.3 (IN), TLS handshake, CERT verify (15):
* TLSv1.3 (IN), TLS handshake, Finished (20):
* TLSv1.3 (OUT), TLS change cipher, Change cipher spec (1):
* TLSv1.3 (OUT), TLS handshake, Finished (20):
* SSL connection using TLSv1.3 / TLS_AES_256_GCM_SHA384
* ALPN: server accepted h2
* Server certificate:
*   subject: CN=wellness.local
*   start date: Feb 10 14:39:16 2026 GMT
*   expire date: May 11 14:39:16 2026 GMT
*   issuer: CN=wellness-ca
*   SSL certificate verify result: unable to get local issuer certificate (20), continuing anyway.
```



Latest Entries

First entry — Initial test entry
Docker ready — Database initialized via init.sql
Ejemplo en Español — Este es un ejemplo en Castellano

Add New Entry

Title

Description